

# The AI Brief.

A weekly intelligence briefing for executives.

|                                   |                                         |                                         |                                       |                                      |
|-----------------------------------|-----------------------------------------|-----------------------------------------|---------------------------------------|--------------------------------------|
| \$145B                            | 80yr                                    | \$10.9B                                 | 11K+                                  | 60%                                  |
| META AI CAPEX 2026 INFRASTRUCTURE | MATH PROBLEM SOLVED OPENAI AUTONOMOUSLY | ANTHROPIC Q2 REV FIRST OPERATING PROFIT | TECH JOBS CUT META + INTUIT THIS WEEK | AI SELF-IMPROVES PROBABILITY BY 2028 |

THIS WEEK — WEEK OF MAY 18 – 24, 2026

- 01 OpenAI's reasoning model autonomously disproves an 80-year-old geometry conjecture, called a milestone in AI mathematics.
- 02 Trump kills the White House AI executive order hours before signing after direct calls from Musk, Zuckerberg, and Sacks.
- 03 Meta fires 8,000 workers while committing \$145B to AI infrastructure; leaked audio reveals employee work trained AI models.
- 04 GitHub confirms 3,800 internal repositories stolen via an 18-minute poisoned VS Code extension window.
- 05 Anthropic reports first operating profit and \$10.9B Q2 revenue; OpenAI files for IPO at a ~\$1.75T valuation.
- 06 Andrej Karpathy joins Anthropic to lead a team accelerating Claude pre-training research using AI.
- 07 Anthropic acquires SDK builder Stainless and closes a \$200M partnership with the Gates Foundation.

# Executive Summary

THE THREE FORCES SHAPING AI THIS WEEK

## THE AI PROOF MOMENT

OpenAI's reasoning model solved an 80-year geometry problem with a 125-page autonomous proof, earning review by Fields medalist Tim Gowers. Anthropic co-founder Jack Clark simultaneously placed 60%+ probability on AI training its own successor by 2028 — now in official research agendas, not theoretical speculation.

**$n^{1.014}$**

ERDOS CONJECTURE IMPROVEMENT

## THE GOVERNANCE GAP

Meta committed \$145B to AI while cutting 8,000 jobs; Anthropic hit its first profit the same week OpenAI filed a \$1.75T IPO. Trump killed the AI safety executive order after direct calls from Musk and Zuckerberg — the starkest governance rollback in the industry's history.

**\$145B**

META'S 2026 AI INFRASTRUCTURE

## THE SUPPLY CHAIN MOMENT

GitHub lost 3,800 internal repos to a worm embedded in a trusted VS Code extension for just 18 minutes; OpenAI and Mistral were also compromised. The attack surface is now standard developer tooling — npm packages, VS Code extensions, CI/CD pipelines — not perimeter networks.

**18 min**

SUPPLY CHAIN BREACH WINDOW

# OpenAI's model autonomously cracks an 80-year math problem

A reasoning model produced a 125-page autonomous proof disproving the Erdos unit distance conjecture — the first time AI has independently solved a major named open math problem.

On May 20, OpenAI announced that an internal general-purpose reasoning model had autonomously disproved the unit distance conjecture, a problem in discrete geometry posed by Paul Erdos in 1946. The model received the problem statement and produced a 125-page proof connecting elementary geometry to algebraic number theory structures mathematicians in combinatorial geometry had never thought to apply here.

The proof method matters more than the result. Rather than iterating on known grid arrangements, the model applied the Golod-Shafarevich criterion from 1964 number theory — a tool from an entirely different branch of mathematics — to show an infinite family of point configurations can produce more unit-distance pairs than any square grid. Princeton's Will Sawin confirmed: the best configurations now scale as  $n^{1.014}$ .

What separates this from previous AI math milestones is the absence of human scaffolding. The model was not trained for this specific problem, did not retrieve an existing solution, and was not guided step-by-step. Fields medalist Tim Gowers reviewed the work and called it 'a milestone in AI mathematics.' Princeton's Noga Alon described it as 'an outstanding achievement' applying tools that surprised the geometry community.

## WHY IT MATTERS

Update research and R&D leadership on what autonomous proof generation means for your scientific function. The barrier between AI as a tool and AI as an independent scientific contributor has shifted materially this week.

---

## Top Stories

---

02 POLICY

4 MIN READ

### Trump scraps AI oversight order after tech CEO intervention

The voluntary 90-day pre-launch review framework was cancelled hours before signing after Musk, Zuckerberg, and Sacks called the president directly.

- WHY IT MATTERS: The federal AI governance vacuum is now official. Build compliance programs around state and international obligations; federal preemption of state AI law appears off the table through 2026.

03 CYBERSECURITY

3 MIN READ

### TeamPCP worm steals 3,800 GitHub repos in 18-minute window

A trojanized VS Code extension was live for 18 minutes on May 18, harvesting credentials from GitHub, OpenAI, Mistral, and the European Commission.

- WHY IT MATTERS: Prioritize signed builds and hardware-backed credential storage across all AI dev toolchains. The breach timeline is now 18 minutes — standard endpoint security cannot respond in that window.

04 ENTERPRISE SHIFT

3 MIN READ

### Meta fires 8,000 and bets \$145B on AI in the same week

Meta reported record \$56.3B revenue the same week 8,000 employees were notified of termination, with leaked audio revealing their work had trained AI models.

- WHY IT MATTERS: Map which roles in your organization are within 18 months of AI displacement. Meta and Intuit's restructuring calculus — costs accrue immediately, retraining does not — will define the political economy of AI through 2027.

05 MARKET STRUCTURE

3 MIN READ

### Anthropic posts first profit as OpenAI confidentially files for IPO

Anthropic reported \$10.9B Q2 revenue and first operating profit the same week OpenAI filed a confidential IPO S-1 targeting a ~\$1.75T late-June listing.

- WHY IT MATTERS: Two frontier lab IPOs in six months will accelerate roadmap pressure. Evaluate both Claude and OpenAI stacks against your workflows now; lock enterprise commitments before public-market incentives realign product priorities.

# Frontier Labs

## Anthropic

FIRST PROFIT · \$900B TALKS

- Andrej Karpathy joins to lead a Claude-accelerated pre-training research team under Nick Joseph.
- Acquires Stainless — the SDK team behind OpenAI's developer tooling — to upgrade Claude API libraries.
- \$200M Gates Foundation partnership over four years for AI in healthcare, education, and agriculture.

## Google DeepMind

I/O 2026 · GEMINI CREATIVE ERA

- Gemini integrates Adobe Firefly, Canva Magic Layers, and CapCut — full creative workflows in one conversation.
- Gemini Spark, a 24/7 personal AI agent, launches for AI Ultra (\$100/mo) subscribers in the US.
- Android XR smart glasses announced with Gentle Monster and Warby Parker; audio model launches fall 2026.

## xAI

FOLDED INTO SPACEX AI · GROK 5

- Grok 4.3 in expanded rollout at \$1.25 / \$2.50 per million tokens — competitive pricing against GPT-5.5.
- Grok Build coding agent moves to wider SuperGrok Heavy beta; direct challenge to Codex and Claude Code.
- Structural rebranding continues: xAI being folded into SpaceX, renamed SpaceX AI.
- Grok 5 expected Q2 2026; Anthropic signed compute deal with SpaceX for Claude model training.

## OpenAI

\$25B ARR · ~\$1.75T IPO TARGET

- Reasoning model solves the 80-year Erdos geometry conjecture with a 125-page autonomous proof.
- Files confidential IPO S-1 targeting ~\$1.75T valuation and a late-June 2026 listing.
- Codex extended to ChatGPT mobile app — developers supervise AI coding agents from a smartphone.
- Integrates SynthID watermarking into all

## Meta AI

\$145B CAPEX · 8,000 JOBS CUT

- Commits \$125–145B to AI in 2026 — nearly double 2025 — while cutting 8,000 jobs (10% of workforce).
- Leaked audio: Zuckerberg admits tracking employee coding sessions to train Meta's AI models.
- Record \$56.3B quarterly revenue and \$26.8B net income in Q1; stock rises on layoff announcement.
- 7,000 workers redirected into new AI teams: Applied AI Engineering and Agent Transformation XFN.

# Model Benchmarks

FRONTIER MODEL SHOWDOWN — CODING, REASONING, AGENTIC TASKS, MAY 2026

| BENCHMARK                    | CLAUDE MYTHOS | GPT-5.5 |
|------------------------------|---------------|---------|
| SWE-bench Pro · coding       | 77.8%         | 58.6%   |
| GPQA Diamond · reasoning     | 94.6%         | ~85%    |
| HLE no tools · science       | 56.8%         | 41.4%   |
| OSWorld · desktop agent      | 79.6%         | 78.7%   |
| BrowseComp · web research    | 86.9%         | 90.1%   |
| Terminal-Bench 2.0           | ~78%          | 82.7%   |
| CyberGym · security          | 83.0%         | 81.8%   |
| GraphWalks BFS · 256K-1M ctx | 80.0%         | ~20%    |

Mythos remains invitation-only via Project Glasswing. GPT-5.5 is live in ChatGPT, Codex, and the API. Gemini 3.5 Flash launched this week at 4x output speed and ~33% of frontier cost.

# Developer Tools

|                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |        |                                                                                                                                                                                                                                                                                                                                                                                                                                                               |        |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------|
| <b>Anthropic Claude Code</b><br>– SDK builder Stainless acquired to improve Python, TypeScript, Go, Ruby, and Java Claude API libraries.<br>– Supply chain alert: rotate credentials immediately if Nx Console VS Code extension was installed<br>– New Managed Agents API lets a single call spin up an agent with persistent state across third-party services.<br>– MCP support for Canva, Instacart, and OpenTable rolling out to AI Ultra subscribers in coming weeks. | MAY 18 | <b>OpenAI Codex</b><br>– Codex extended to ChatGPT iOS and Android apps — run and supervise coding agents from a smartphone.<br>– OpenAI macOS app code-signing certificate revoked June 12 after supply chain breach; reinstall required.<br>– Cloud agent dev environments and Microsoft Teams integration for async AI coding workflows now live.<br>– Build-in-Parallel runs multiple AI sub-agents concurrently; priced at \$0.50 / \$2.50 per M tokens. | MAY 20 |
| <b>Gemini Managed Agents</b>                                                                                                                                                                                                                                                                                                                                                                                                                                                | MAY 19 | <b>Cursor Composer 2.5</b>                                                                                                                                                                                                                                                                                                                                                                                                                                    | MAY 18 |

# Research Highlights

## Erdos Unit Distance Disproof

OPENAI RESEARCH · OPENAI.COM

A reasoning model autonomously produced a 125-page proof disproving a 1946 geometry conjecture, connecting elementary geometry to Golod-Shafarevich class field towers. Confirmed by Fields medalist Tim Gowers; Princeton's Sawin quantified:  $n^{1.014}$ .

## AI-Accelerated AI: Early Signs

ANTHROPIC INSTITUTE · MAY 2026

Anthropic published a five-page agenda documenting early signs of AI contributing to its own R&D. Jack Clark placed 60%+ probability on AI training its own successor by 2028 — the first mainstream publication to use 'intelligence explosion' with a specific timeframe.

## Neuro-Symbolic Robotics

TUFTS UNIVERSITY · ICRA VIENNA

A neuro-symbolic system combining neural networks with symbolic reasoning achieved 95% success in robotic task tests versus 34% for standard models, training in 34 minutes compared to 36+ hours. Results presented at ICRA.

## MOSS: Agent Self-Evolution

ARXIV · MAY 2026

MOSS introduces source-level code rewriting as a mechanism for autonomous agent self-improvement, allowing agents to modify their own operational logic between tasks without external intervention.

## Gated DeltaNet-2 Attention

HATAMIZADEH ET AL. · ARXIV 2026

Introduces a linear attention architecture that decouples erase and write operations, achieving performance competitive with full self-attention at significantly lower computational cost for long-context inference.

## MotiMotion: Motion Video Gen

ARXIV · MAY 21, 2026

MotiMotion enhances motion-controlled video generation by integrating visual reasoning and confidence-aware control, enabling models to reason about physical plausibility before generating motion trajectories.

---

# AI Index — Key Findings

---

## AUTOMATION WAVE

---

Meta (8,000 jobs), Intuit (3,000 jobs), and LinkedIn cut headcount this week while simultaneously raising AI infrastructure spending. More than 11,000 positions were eliminated in seven days; productivity gains are accruing to shareholders, not workers.

## AI SCIENTIFIC AGENCY

---

OpenAI's math proof represents a category shift: a model received a problem statement and produced a 125-page novel proof entirely autonomously — no targeted training, no step-by-step human guidance, and no prior solution to retrieve.

## IPO PIPELINE

---

Anthropic (targeting October 2026) and OpenAI (targeting late June 2026) are both preparing public listings in the same window. A combined implied valuation of ~\$2.65T would make this the largest simultaneous IPO pair in market history.

## DEVELOPER RISK SURFACE

---

The TeamPCP supply chain campaign hit 170+ npm packages, two PyPI libraries, and a 2.2M-install VS Code extension — compromising GitHub, OpenAI, Mistral, and the European Commission from standard developer tools with no perimeter breach required.

## CREATIVE AI SHIFT

---

Google's Gemini now orchestrates Adobe Firefly, Canva, and CapCut inside a single conversation. The creative workflow previously requiring four applications and three tabs is now one prompt chain — reshaping the economics of solo creators and marketing teams.



# VC & Funding

|                             |                     |                          |                           |
|-----------------------------|---------------------|--------------------------|---------------------------|
| <b>\$145B</b>               | <b>\$950M</b>       | <b>\$230M</b>            | <b>\$200M</b>             |
| META AI CAPEX 2026          | SIERRA SERIES ROUND | PARALLEL AGENT INFRA     | GATES X ANTHROPIC         |
| NEARLY DOUBLE 2025 SPENDING | \$15.8B VALUATION   | AI WEB-SEARCH FOR AGENTS | 8-YEAR AI DEPLOYMENT DEAL |

## MEGA ROUNDS

| COMPANY           | ROUND            | AMOUNT  | LEAD              | VALUATION          | NOTE                        |
|-------------------|------------------|---------|-------------------|--------------------|-----------------------------|
| Sierra            | Series E         | \$950M  | Tiger Global · GV | \$15.8B post-money | Enterprise AI custome...    |
| Parallel          | Series B         | \$230M  | Undisclosed       | —                  | AI-native web search f...   |
| Anthropic         | Pre-IPO round    | \$30B+  | GIC · Coatue      | \$900B (talks)     | First operating profit Q... |
| OpenAI            | Confidential ... | —       | Public markets    | ~\$1.75T target    | Confidential S-1 filed ...  |
| Gates x Anthropic | Strategic        | \$200M  | Gates Foundation  | —                  | 4-year deal for AI in he... |
| Profluent x Lilly | Strategic        | \$2.25B | Eli Lilly         | —                  | AI protein design for g...  |

## ENTERPRISE MOVES

- Meta** Commits \$125–145B to AI in 2026 — nearly double 2025 — while cutting 8,000 jobs, redirecting to Applied AI Engineering and Agent Transformation.
- Intuit** Cuts 3,000 jobs (8% of workforce) to redirect \$500M+ in annualized savings toward AI agents, including AI-powered QuickBooks and TurboTax automation.
- Anthropic** Acquires Stainless, hires Karpathy to lead pre-training acceleration, and reports first quarterly operating profit of \$559M on \$10.9B Q2 revenue.
- Google DeepMind** Integrates Adobe Firefly, Canva, and CapCut into Gemini — full professional creative workflows in one AI conversation for Ultra (\$100/mo) subscribers.
- OpenAI** Extends Codex to ChatGPT mobile; files confidential IPO S-1 at ~\$1.75T; revokes macOS certificate June 12 post-supply-chain breach.

## Robotics & Science

### Tufts Neuro-Symbolic Robots

95% task success vs 34% for standard models; trains in 34 minutes vs 36+ hours; to be presented at ICRA Vienna.

### NVIDIA x Cadence Sim-to-Real

Expanded partnership combines multiphysics simulation with Isaac robotics and Cosmos open-world models to close training-deployment gap.

### Figure AI Series F

\$1B raised; Figure 03 entering mass-manufacturing validation for automotive and electronics assembly lines.

### Unitree G1 Humanoid

Priced at \$16,000 and shipping to commercial buyers — first sub-\$20K bipedal robot available for enterprise evaluation.

## Regulatory Tracker

### United States

HIGH

White House AI EO killed: Trump kills voluntary 90-day pre-launch model review framework hours before signing, after direct calls from Musk, Zuckerberg, and David Sacks.

### United States

MEDIUM

Take It Down Act live: Effective May 19: platforms must remove non-consensual intimate images — including AI deepfakes — within 48 hours of a victim's notice.

### European Union

HIGH

EU AI Act — Aug. activation: Technical annex comment period open to June 30; systemic-risk evaluation guidance published for GPAI providers above compute thresholds.

### United Kingdom

MEDIUM

AI SI pre-deploy methodology: UK AI Safety Institute publishes methodology for pre-deployment model evaluation, signaling enforcement direction for forthcoming statutory regime.

### US (State level)

MEDIUM-HIGH

State AI bills accelerating: Washington, Florida, Virginia, and Utah advance AI legislation on consumer rights and transparency as the federal governance vacuum deepens.

## Thought Leaders

"My prediction is by the end of 2028, it's more likely than not that we have an AI system where you would be able to say to it: Make a better version of yourself."

**Jack Clark**

"This is a milestone in AI mathematics."

**Tim Gowers**

CAMBRIDGE / FIELDS MEDAL

"The AI models learn from watching really smart people do things. No human is looking at or watching what people are doing on their computers."

**Mark Zuckerberg**

"I think the next few years at the frontier of LLMs will be especially formative. I am very excited to join the team here and get back to R&D."

**Andrej Karpathy**

ANTHROPIC

Community Pulse

X / AI COMMUNITY

OpenAI's math proof ignited debate: is autonomous proof-writing genuine mathematical understanding, or sophisticated pattern matching at R/MACHINELEARNING

Jack Clark's 60% recursive self-improvement estimate sparked debate: if AI training AI is already happening deliberately, is the safety community's R/DEVOPS / HACKERNEWS

TeamPCP supply chain attack dominated developer forums: 18-minute breach window in a trusted extension raises calls for signed packages and hardware-backed developer LINKEDIN / TECH

Meta training AI on employees' work before mass layoffs reframed the debate: the first step of automation may be using your work to build the system that replaces you.

What to Watch

**June 11–12    OpenAI IPO pricing window**  
OpenAI targets late-June listing; final S-1 detail will clarify AI lab unit economics and governance restructuring for public investors.

**June 12        OpenAI macOS cert revocation**  
macOS ChatGPT becomes untrusted after June 12; a forced reinstall from a freshly signed build is required — end users not yet notified.

**June 30        EU AI Act comment closes**  
Last window to shape GPAI systemic-risk evaluation obligations before they harden; any EU-facing AI deployment needs a filed response.

**June (WWDC)   Apple AI Extensions**  
Whether Apple opens iOS to Anthropic and Google agents could redraw consumer AI distribution for hundreds of millions of users.

**October 2026   Anthropic IPO launch**  
First-ever operating profit and \$10.9B Q2 revenue sets the stage; valuation talks near \$900B; will be the largest AI lab IPO on record.

Sources

Build Fast with AI. AI News Today — May 23, 2026: 12 Biggest Stories. buildfastwithai.com.  
OpenAI. Model disproves discrete geometry conjecture. openai.com/index. May 20, 2026.  
Axios. Why Trump's AI executive order was pulled. May 21, 2026.  
TechCrunch. OpenAI co-founder Andrej Karpathy joins Anthropic's pre-training team. May 19, 2026.  
The Hacker News. GitHub repos breached via poisoned Nx Console VS Code extension. May 20, 2026.  
NPR. Meta slashes 8,000 jobs as it pivots toward AI. May 20, 2026.  
Fortune. The big questions OpenAI's trillion-dollar IPO filing may finally answer. May 22, 2026.  
Time. Anthropic sells Claude's promise while warning about AI's dangers. Jack Clark, May 22, 2026.  
Anthropic. Anthropic acquires Stainless. anthropic.com/news. May 18, 2026.

Glossary

**Supply chain attack**  
A cyber intrusion exploiting trusted software dependencies — packages, extensions — rather than attacking a target's perimeter directly.  
**Unit distance problem**  
A 1946 geometry problem: how many point-pairs can be exactly distance 1 apart. OpenAI's model disproved the conjectured upper bound this week.  
**MCP**  
Model Context Protocol. Open standard for tool calling between AI models and external services; 97M installs across all major AI providers.  
**Agentic system**  
An AI system that plans and executes multi-step tasks autonomously, typically with tool use and persistent state across a session.

**Recursive improvement**  
A process where an AI system improves its own training, potentially producing exponentially more capable successors without human direction.  
**Pre-training**  
The initial large-scale training run giving an AI model its core knowledge. Karpathy joined Anthropic to use Claude to accelerate this stage itself.  
**GPAI**  
General-Purpose AI. EU AI Act category for broad-capability foundation models above compute thresholds, subject to enhanced obligations.