


The AI Brief.

A weekly intelligence briefing for executives.



OPENAI IPO TARGET

\$1.75T

Roadshow opens June 2

ANTHROPIC VALUE

\$1T

Talks near term sheet

ERDOS BOUND

$n^{1.05}$

Second proof verified

EU AI ACT CLOCK

62 d.

Full force in August

MCP INSTALLS

104M

Up from 97M last week

TODAY'S HEADLINES

6 SIGNALS

- 01 OpenAI opens its IPO roadshow June 2, targeting a \$1.75T valuation and a late-June listing on the Nasdaq.
- 02 Anthropic enters term-sheet talks at a \$1T valuation weeks after posting its first operating profit.
- 03 Gowers confirms OpenAI's Erdos proof and a second model run tightens the unit-distance bound to $n^{1.05}$.
- 04 Apple previews third-party AI Extensions for iOS at WWDC, opening Siri to Anthropic and Google models.
- 05 GitHub mandates signed builds and artifact attestation across all repos after the TeamPCP supply-chain worm.
- 06 EU AI Office publishes final systemic-risk guidance as the Act's full applicability nears on August 2.

SECTION ONE

Executive summary.

The themes that frame this week's signals — and the reasoning behind why each belongs on a leadership team's agenda.

PREPARED FOR LEADERSHIP · 8 MIN READ · CURATED BY THE BRIEF DESK

01

THE IPO WINDOW OPENS

Two frontier labs now have public listings in flight inside one quarter. OpenAI's roadshow opens June 2 at a \$1.75T target while Anthropic fields a \$1T term sheet — pricing both labs' unit economics in public for the first time.

\$1.75T

OpenAI IPO valuation target

02

AUTONOMOUS SCIENCE

AI as an independent research contributor moved from claim to record this week. Gowers verified OpenAI's Erdos proof and a second run tightened the bound to $n^{1.05}$, while Anthropic's Karpathy-led team reported its first Claude-accelerated pre-training gains.

$n^{1.05}$

New unit-distance lower bound

03

GOVERNANCE HARDENS

Regulation and security tightened in tandem. The EU published final systemic-risk guidance 62 days before the Act binds, and GitHub mandated signed builds across every repository — making secure delivery an engineering line item, not a policy memo.

62 d.

Until EU AI Act full force

LEAD — MARKET STRUCTURE

6 MIN READ

OpenAI opens its IPO roadshow at a \$1.75T valuation.

OpenAI begins its investor roadshow June 2 ahead of a late-June Nasdaq listing — the first frontier-lab IPO and the largest technology offering on record.

WHY IT MATTERS

Lock enterprise pricing and SLAs with both labs before June pricing. Public-market incentives will push OpenAI and Anthropic toward margin discipline within two quarters — free tiers and soft commitments are the first things to go.

OpenAI filed its final S-1 this week and will open an investor roadshow on June 2, targeting a \$1.75T valuation and a late-June listing on the Nasdaq. The offering would be the largest technology IPO in history, eclipsing the prior record by a wide margin and converting the most-watched private company in AI into a public one inside a single quarter.

The filing discloses unit economics for the first time. OpenAI reported \$25B in annualized revenue against heavy compute and talent costs, a Microsoft partnership restructured to remove exclusivity, and a governance structure that survived the Musk litigation intact. Public investors will price the gap between the company's research ambitions and its current margins.

Anthropic is moving on a parallel track. Fresh off its first operating profit, it has entered term-sheet talks at a \$1T valuation and is weighing an October listing. Two frontier labs pricing in public within months will set the comparable that every downstream AI valuation is measured against.

For buyers, the listings change the incentive structure of the vendors you depend on. Public-market scrutiny will pull both labs toward margin discipline — fewer free tiers, firmer enterprise pricing, faster depreciation of unprofitable features. Lock multi-year terms now, before the roadshow narrative hardens into roadmap.

SECTION TWO · MORE SIGNALS

Stories of the week.

Four more developments that warrant leadership attention this week.

02

SCIENTIFIC DISCOVERY

3 MIN READ

Gowers verifies OpenAI's Erdos proof as a second run tightens the bound.

Fields medalist Tim Gowers completed his review and a second autonomous run improved the unit-distance lower bound to $n^{1.05}$, ending doubts the first proof was a fluke.

WHY IT MATTERS

Brief R&D leadership that autonomous proof generation is now reproducible, not a one-off. Begin scoping which internal research questions a frontier model could attempt this year.

03

CONSUMER PLATFORM

3 MIN READ

Apple previews third-party AI Extensions for iOS at WWDC.

Apple will let users route Siri and system requests to Anthropic and Google models through a new Extensions framework, opening its install base to outside frontier labs.

WHY IT MATTERS

If your product touches iOS, plan for Siri to broker requests to third-party models by fall. Decide now whether you integrate as an Extension provider or defend your own surface.

04

CYBERSECURITY

3 MIN READ

GitHub mandates signed builds and attestation after the TeamPCP worm.

Every repository must adopt build provenance and artifact attestation by July 1 in GitHub's response to the 18-minute supply-chain breach that hit OpenAI, Mistral, and the EU.

WHY IT MATTERS

Budget for build-provenance rollout across your toolchain this quarter. Treat unsigned dependencies and editor extensions as the primary breach vector now, not the network perimeter.

05

POLICY

3 MIN READ

EU publishes final systemic-risk guidance ahead of August activation.

The AI Office released binding evaluation guidance for GPAI providers above the compute threshold; the comment window on the technical annex closes June 30.

WHY IT MATTERS

Any GPAI deployment reaching EU residents needs an evaluation programme described and budgeted before August 2. File comments on the technical annex before June 30 to shape the obligations.

SECTION THREE · FRONTIER MODELS

Labs, in their own words.

Per-lab summary of the week's model and capability developments — the commercial map of the frontier.

Anthropic

\$1T TALKS · FIRST PROFIT

- Enters term-sheet talks at a \$1T valuation weeks after reporting its first operating profit.
- Karpathy's team reports early Claude-accelerated pre-training gains under Nick Joseph.
- Closes the Stainless acquisition and ships rebuilt Python, TypeScript, and Go API libraries.
- Signs a multi-year SpaceX compute deal to secure training capacity for the next Claude generation.

OpenAI

\$25B ARR · \$1.75T IPO TARGET

- Files its final S-1 and opens an investor roadshow June 2 ahead of a late-June Nasdaq listing.
- Second autonomous run tightens the Erdos unit-distance bound to $n^{1.05}$, confirming the proof.
- Restructures the Microsoft partnership to remove exclusivity and broaden cloud relationships.
- Extends Codex supervision to the ChatGPT mobile apps for iOS and Android.

Google DeepMind

I/O FOLLOW-THROUGH · GEMINI

- Gemini 3.5 Flash ships broadly at 4x the output speed of frontier peers at roughly one-third the cost.
- Gemini Spark personal agent opens to AI Ultra subscribers across the US.
- Adds a Managed Agents API: one call spins up a stateful agent across third-party services.
- Creative stack expands as Firefly, Canva, and CapCut integrations reach general availability.

Meta AI

\$145B CAPEX · 8,000 CUT

- Stands up Applied AI Engineering and Agent Transformation teams from 7,000 redirected staff.
- Confirms \$125–145B in 2026 AI infrastructure spend, nearly double the prior year.
- Faces scrutiny over leaked audio describing training on employee work before the layoffs.
- Posts record \$56.3B quarterly revenue as the stock climbs on the restructuring.

xAI

FOLDED INTO SPACEX AI · GROK 5

- Grok 5 enters internal testing with a public beta still expected this quarter.
- Grok Build coding agent widens its SuperGrok Heavy beta against Codex and Claude Code.
- Rebranding to SpaceX AI proceeds, consolidating governance and compute under one entity.
- Holds aggressive pricing at \$1.25 / \$2.50 per million tokens to undercut larger rivals.

SECTION FOUR · BENCHMARKS & DEVELOPER TOOLS

The frontier, head-to-head.

Claude Mythos vs GPT-5.5 — head-to-head, May 2026

BENCHMARK	CLAUDE MYTHOS	GPT-5.5
SWE-bench Pro · coding	77.8%	58.6%
GPQA Diamond · reasoning	94.6%	~85%
HLE no tools · science	56.8%	41.4%
OSWorld · desktop agent	79.6%	78.7%
BrowseComp · web research	86.9%	90.1%
Terminal-Bench 2.0	~78%	82.7%
CyberGym · security	83.0%	81.8%
GraphWalks BFS · long ctx	80.0%	~20%

Mythos remains invitation-only via Project Glasswing; GPT-5.5 is live in ChatGPT, Codex, and the API. Gemini 3.5 Flash now ships at roughly one-third of frontier cost.

Developer tools — week-in-progress

Claude Code

UPD. MAY 27

- Rebuilt API libraries from the Stainless team land for Python, TypeScript, Go, Ruby, and Java.
- Build-provenance checks added by default; unsigned dependencies are flagged before each run.

OpenAI Codex

UPD. MAY 26

- Mobile supervision reaches general availability on the ChatGPT iOS and Android apps.
- Adds signed-build enforcement after the macOS certificate revocation earlier this month.

Cursor Composer 2.5

UPD. MAY 25

- Build-in-Parallel runs multiple sub-agents concurrently at \$0.50 / \$2.50 per M tokens.
- Cloud agent environments and Microsoft Teams integration reach all paid tiers.

Gemini Managed Agents

UPD. MAY 28

- New API spins up a stateful agent across third-party services in a single call.
- MCP support for Canva, Instacart, and OpenTable rolls out to AI Ultra subscribers.

SECTION FIVE · RESEARCH

Papers and findings.

What the academic literature and the major AI indices added this week.

Notable research

Erdos Unit Distance, Verified

OPENAI RESEARCH · OPENAI.COM

Tim Gowers completed peer review of the autonomous 125-page proof and a second model run tightened the lower bound to $n^{1.05}$, establishing the result as reproducible rather than a single fortunate trajectory.

Claude-Accelerated Pre-training

ANTHROPIC · TECHNICAL REPORT

Anthropic's Karpathy-led team reports the first measurable pre-training efficiency gains from using Claude to optimize its own training pipeline — an early data point for recursive R&D acceleration.

Neuro-Symbolic Robotics

TUFTS UNIVERSITY · ICRA VIENNA

A hybrid system pairing neural perception with symbolic reasoning reached 95% task success versus 34% for standard models, training in 34 minutes against 36+ hours, presented at ICRA.

MOSS: Agent Self-Evolution

ARXIV · MAY 2026

MOSS introduces source-level code rewriting as a mechanism for autonomous self-improvement, letting an agent modify its own operational logic between tasks without external intervention.

Gated DeltaNet-2 Attention

HATAMIZADEH ET AL. · ARXIV 2026

A linear-attention architecture that decouples erase and write operations, matching full self-attention quality at significantly lower compute for long-context inference.

Warmth Tax in LLMs

NATURE

Models tuned to be warm and agreeable are measurably more likely to give incorrect answers and reinforce false beliefs, quantifying a trade-off between likability and reliability.

Stanford HAI Index — key findings

IPO PIPELINE	OpenAI (late June) and Anthropic (targeting October) are both preparing public listings in the same window. A combined implied valuation near \$2.75T would be the largest simultaneous frontier-lab debut on record.
AUTONOMOUS SCIENCE	Verification of OpenAI's Erdos proof and a tighter second result mark AI's shift from research tool to independent contributor — a problem statement in, a novel 125-page proof out, no human scaffolding.
CAPABILITY CEILING	SWE-bench Verified climbed from 60% to near 100% in a single year. Frontier models now meet or exceed human baselines on PhD-level science, multimodal reasoning, and competition mathematics.
GOVERNANCE GAP	The Foundation Model Transparency Index fell from 58 to 40 out of 100 in a year even as capability rose — the most capable models are now the least transparent about training and evaluation.
WORKFORCE	More than 11,000 tech roles were cut in a single week at firms simultaneously raising AI spending. Entry-level developer hiring for ages 22–25 is down roughly 20% since 2024.

SECTION SIX · CAPITAL

The funding landscape.

OPENAI IPO TARGET	ANTHROPIC VALUATION	SIERRA SERIES E	PROFLUENT × LILLY
\$1.75T	\$1T	\$950M	\$2.25B
Largest tech offering on record	Term-sheet talks underway	\$15.8B post-money	AI protein design pact

Mega rounds & strategic bets

COMPANY	ROUND	AMOUNT	LEAD	VALUATION
OpenAI Final S-1 filed; roadshow opens June 2 on the Nasdaq.	Pre-IPO / public	—	Public markets	\$1.75T target
Anthropic Term sheet under negotiation after first operating profit.	Pre-IPO round	\$30B+	GIC · Coatue	\$1T (talks)
Sierra Enterprise AI agents for customer operations.	Series E	\$950M	Tiger Global · GV	\$15.8B post-money
Parallel AI-native web search built for autonomous agents.	Series B	\$230M	Undisclosed	—
Profluent AI protein design for large-gene insertion therapies.	Strategic	\$2.25B	Eli Lilly	—
Figure AI Figure 03 enters mass-manufacturing validation.	Series F	\$1B	—	\$45B (talks)
Sereact Embodied AI for warehouse manipulation tasks.	Series B	\$110M	—	—

SECTION SEVEN · INDUSTRY · ROBOTICS · POLICY

Where the rules and runways meet.

Enterprise & industry moves

JPMorgan Chase	Scales its \$19.8B tech budget toward agentic workflows, with 2,000 AI staff scanning \$10T in daily transactions for a projected \$2.5B in annual value.
Microsoft	Confirms a \$10B four-year Japan commitment for data centers with SoftBank and pledges to train more than one million engineers by 2030.
Novo Nordisk	Expands its OpenAI deployment across discovery, trials, and manufacturing, targeting full enterprise integration by the end of 2026.
Intuit	Redirects more than \$500M in savings from 3,000 job cuts into AI agents for QuickBooks and TurboTax automation.
Anthropic	Closes the Stainless acquisition and a \$200M, four-year Gates Foundation partnership spanning health, education, and agriculture.
AWS AgentCore	Lets AI agents make USDC micropayments via the x402 protocol through a new Coinbase and Stripe partnership.

Robotics & science

Tufts Neuro-Symbolic Robots 95% task success versus 34% for standard models; trains in 34 minutes against 36+ hours, presented at ICRA Vienna.	NVIDIA × Cadence Expanded partnership pairs multiphysics simulation with Isaac robotics and Cosmos world models to close the sim-to-real gap.
Figure 03 First humanoid line to enter mass-manufacturing validation for automotive and electronics assembly	Unitree G1 Sub-\$20K bipedal robot ships to commercial buyers at \$16,000 for enterprise evaluation

Regulatory & policy

EUROPEAN UNION	AI Act — full force August 2 Final systemic-risk evaluation guidance published for GPAI providers above the compute threshold; annex comment window closes June 30.	IMPACT High
UNITED STATES	Federal AI order rescinded The voluntary pre-launch review framework remains cancelled; compliance shifts to state and international regimes through 2026.	IMPACT Medium-High
UNITED STATES	Take It Down Act in force Platforms must remove non-consensual intimate images, including AI deepfakes, within 48 hours of a victim's notice.	IMPACT Medium
UNITED KINGDOM	AISI pre-deployment methodology The AI Safety Institute's published evaluation methodology signals the enforcement direction for a forthcoming statutory regime.	IMPACT Medium

SECTION EIGHT · VOICES · PULSE · WATCH

What the field is saying.

Thought leader pulse

“Going public is not the finish line. It is the capital base for the largest research program in the history of the field.”

— SAM ALTMAN · OPENAI

On the IPO roadshow opening June 2.

“I have now checked the argument in full. This is a milestone in AI mathematics, not a curiosity.”

— TIM GOWERS · CAMBRIDGE / FIELDS MEDAL

On verifying the Erdos unit-distance proof.

“Profitability buys us independence. It lets us set the safety agenda instead of asking permission to have one.”

— DARIO AMODEI · ANTHROPIC

After Anthropic's first operating profit.

“The early pre-training gains are small but real. We are watching a model help build the next model, and that is the whole point.”

— ANDREJ KARPATHY · ANTHROPIC

On the first Claude-accelerated training results.

Community pulse

X / AI COMMUNITY	OpenAI's \$1.75T target reignited the bubble debate: are public markets about to price frontier research correctly, or repeat the dot-com overshoot at a larger scale?
R/MACHINELEARNING	Karpathy's pre-training results drew measured optimism — commenters want the ablations and baselines before calling it the start of recursive self-improvement.
R/DEVOPS / HACKERNEWS	GitHub's signed-build mandate dominated threads, with debate over whether attestation actually closes the gap a trusted extension opened in 18 minutes.
LINKEDIN / TECH	Apple's rumored AI Extensions split opinion: a distribution windfall for Anthropic and Google, or one more tax-and-database layer for developers.

What to watch

JUNE 2 OpenAI roadshow opens The investor roadshow begins; final pricing detail will clarify frontier-lab unit economics for public markets.	JUNE (WWDC) Apple AI Extensions Whether Apple opens iOS to Anthropic and Google models could redraw consumer AI distribution for hundreds of millions.	JUNE 30 EU comment window closes Last chance to shape GPAI systemic-risk evaluation obligations before they harden ahead of August activation.	JULY 1 GitHub attestation deadline Signed-build and artifact-attestation requirements take effect across all repositories.	AUGUST 2 EU AI Act full force Full applicability lands; any non-compliant GPAI deployment reaching EU residents enters enforcement territory.
--	--	--	--	---

END MATTER

Sources & glossary.

The reading list, the working vocabulary, and where this issue ends.

Sources

01	Stanford HAI. 2026 AI Index Report. May 2026.
02	OpenAI. Form S-1 registration statement and IPO roadshow materials. May 2026.
03	OpenAI. Second model run improves discrete-geometry bound. openai.com/index . May 2026.
04	Anthropic. First operating profit and \$1T valuation talks. anthropic.com/news . May 2026.
05	EU AI Office. Final Guidance on Systemic-Risk Evaluations. Brussels, May 2026.
06	GitHub. Mandatory build provenance and artifact attestation. github.blog . May 2026.
07	Apple. WWDC 2026 preview briefings on AI Extensions for iOS. May 2026.
08	PitchBook / CB Insights. Q2 2026 Venture Capital Tracker.
09	Company filings, press releases, and primary-source interviews under Chatham House rules.

Glossary

Roadshow. A series of investor presentations a company gives before an IPO to build demand and inform final pricing.	S-1. The registration statement a US company files with the SEC before going public, disclosing financials and risks.
GPAL. General-Purpose AI. EU AI Act category for broad-capability foundation models above a compute threshold.	Build provenance. Cryptographically verifiable metadata recording how and where a software artifact was built.
MCP. Model Context Protocol. Open standard for tool calling and context exchange between models and applications.	Agentic system. An AI system that plans and executes multi-step tasks autonomously, typically with tool use and persistent state.
Recursive R&D. Using an AI system to improve its own training or successor, compounding capability with less human direction.	

◆ End of Vol. 02 · Issue 23 · Week of May 25 – 31, 2026 ◆